

▼ Libraries

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import tensorflow as tf
from tensorflow import keras
from tensorflow.keras import layers

from keras.applications.densenet import DenseNet121
from keras.applications.xception import Xception
from keras.applications.densenet import preprocess_input
from keras.optimizers import Adam, SGD, Adamax
from keras.models import Model, load_model
from keras.layers import *
from sklearn.model_selection import train_test_split
from keras.callbacks import *

from tensorflow.keras.callbacks import EarlyStopping
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, Embedding, LSTM, Conv1D, MaxPool1D
from sklearn.model_selection import train_test_split
from sklearn.metrics import classification_report, accuracy_score
from tensorflow.keras import regularizers
```

▼ Train Test Split

```
from google.colab import drive
drive.mount('/content/drive', force_remount=True)

Mounted at /content/drive

image_size = (224, 224)
batch_size = 32

train_ds, val_ds = tf.keras.utils.image_dataset_from_directory(
    "/content/drive/MyDrive/BM1000",
    validation_split=0.25,
    subset="both",
    seed=1337,
    image_size=image_size,
    batch_size=batch_size,
)

Found 5850 files belonging to 7 classes.
Using 4388 files for training.
Using 1462 files for validation.

class_names = train_ds.class_names
print(class_names)
num_classes = len(class_names)
print(num_classes)

['BAS', 'EOS', 'HAC', 'LYT', 'MON', 'NGB', 'NGS']
```

▼ Data Augmentation

```
data_augmentation = keras.Sequential(
    [
        layers.RandomFlip("horizontal",
                           input_shape=(224,
                                       224,
                                       3)),
    ])
```

```

        layers.RandomRotation(0.2),
        layers.RandomZoom(0.2),
        layers.RandomContrast(0.2),
        layers.RandomBrightness(0.2)
    ]
)

train_ds = train_ds.map(
    lambda x, y: (data_augmentation(x, training=True), y))

```

▼ Configure the dataset for performance

Make sure to use buffered prefetching, so you can yield data from disk without having I/O become blocking. These are two important methods you should use when loading data:

`Dataset.cache` keeps the images in memory after they're loaded off disk during the first epoch. This will ensure the dataset does not become a bottleneck while training your model. If your dataset is too large to fit into memory, you can also use this method to create a performant on-disk cache.

`Dataset.prefetch` overlaps data preprocessing and model execution while training.

```

AUTOTUNE = tf.data.AUTOTUNE

train_ds = train_ds.cache().shuffle(1000).prefetch(buffer_size=AUTOTUNE)
val_ds = val_ds.cache().prefetch(buffer_size=AUTOTUNE)

```

▼ Custom Callback

```

import time
### Define a class for custom callback
class MyCallback(keras.callbacks.Callback):
    def __init__(self, model, base_model, patience, stop_patience, threshold, factor, batches, initial_epoch, epochs, ask_epoch):
        super(MyCallback, self).__init__()
        self.model = model
        self.base_model = base_model
        self.patience = patience # specifies how many epochs without improvement before learning rate is adjusted
        self.stop_patience = stop_patience # specifies how many times to adjust lr without improvement to stop training
        self.threshold = threshold # specifies training accuracy threshold when lr will be adjusted based on validation loss
        self.factor = factor # factor by which to reduce the learning rate
        self.batches = batches # number of training batch to run per epoch
        self.initial_epoch = initial_epoch
        self.epochs = epochs
        self.ask_epoch = ask_epoch
        self.ask_epoch_initial = ask_epoch # save this value to restore if restarting training
        # callback variables
        self.count = 0 # how many times lr has been reduced without improvement
        self.stop_count = 0
        self.best_epoch = 1 # epoch with the lowest loss
        self.initial_lr = float(tf.keras.backend.get_value(model.optimizer.lr)) # get the initial learning rate and save it
        self.highest_tracc = 0.0 # set highest training accuracy to 0 initially
        self.lowest_vloss = np.inf # set lowest validation loss to infinity initially
        self.best_weights = self.model.get_weights() # set best weights to model's initial weights
        self.initial_weights = self.model.get_weights() # save initial weights if they have to get restored

    # Define a function that will run when train begins
    def on_train_begin(self, logs= None):
        msg = '{0:^8s}{1:^10s}{2:^9s}{3:^9s}{4:^9s}{5:^9s}{6:^9s}{7:^10s}{8:^10s}{9:^8s}'.format('Epoch', 'Loss', 'Accuracy', 'V_loss', 'V_acc',
        print(msg)
        self.start_time = time.time()

    def on_train_end(self, logs= None):
        stop_time = time.time()
        tr_duration = stop_time - self.start_time
        hours = tr_duration // 3600
        minutes = (tr_duration - (hours * 3600)) // 60
        seconds = tr_duration - ((hours * 3600) + (minutes * 60))
        msg = f'training elapsed time was {str(hours)} hours, {minutes:.1f} minutes, {seconds:.2f} seconds'
        print(msg)
        self.model.set_weights(self.best_weights) # set the weights of the model to the best weights

```

```

def on_train_batch_end(self, batch, logs= None):
    acc = logs.get('accuracy') * 100 # get batch accuracy
    loss = logs.get('loss')
    msg = '{0:20s}processing batch {1:} of {2:5s}-  accuracy= {3:5.3f} -  loss: {4:8.5f}'.format(' ', str(batch), str(self.batches),
    print(msg, '\r', end= '') # prints over on the same line to show running batch count

def on_epoch_begin(self, epoch, logs= None):
    self.ep_start = time.time()

# Define method runs on the end of each epoch
def on_epoch_end(self, epoch, logs= None):
    ep_end = time.time()
    duration = ep_end - self.ep_start

    lr = float(tf.keras.backend.get_value(self.model.optimizer.lr)) # get the current learning rate
    current_lr = lr
    acc = logs.get('accuracy') # get training accuracy
    v_acc = logs.get('val_accuracy') # get validation accuracy
    loss = logs.get('loss') # get training loss for this epoch
    v_loss = logs.get('val_loss') # get the validation loss for this epoch

    if acc < self.threshold: # if training accuracy is below threshold adjust lr based on training accuracy
        monitor = 'accuracy'
        if epoch == 0:
            pimprov = 0.0
        else:
            pimprov = (acc - self.highest_tracc ) * 100 / self.highest_tracc # define improvement of model progress

        if acc > self.highest_tracc: # training accuracy improved in the epoch
            self.highest_tracc = acc # set new highest training accuracy
            self.best_weights = self.model.get_weights() # training accuracy improved so save the weights
            self.count = 0 # set count to 0 since training accuracy improved
            self.stop_count = 0 # set stop counter to 0
            if v_loss < self.lowest_vloss:
                self.lowest_vloss = v_loss
            self.best_epoch = epoch + 1 # set the value of best epoch for this epoch

        else:
            # training accuracy did not improve check if this has happened for patience number of epochs
            # if so adjust learning rate
            if self.count >= self.patience - 1: # lr should be adjusted
                lr = lr * self.factor # adjust the learning by factor
                tf.keras.backend.set_value(self.model.optimizer.lr, lr) # set the learning rate in the optimizer
                self.count = 0 # reset the count to 0
                self.stop_count = self.stop_count + 1 # count the number of consecutive lr adjustments
                self.count = 0 # reset counter
                if v_loss < self.lowest_vloss:
                    self.lowest_vloss = v_loss
            else:
                self.count = self.count + 1 # increment patience counter

    else: # training accuracy is above threshold so adjust learning rate based on validation loss
        monitor = 'val_loss'
        if epoch == 0:
            pimprov = 0.0
        else:
            pimprov = (self.lowest_vloss - v_loss ) * 100 / self.lowest_vloss
        if v_loss < self.lowest_vloss: # check if the validation loss improved
            self.lowest_vloss = v_loss # replace lowest validation loss with new validation loss
            self.best_weights = self.model.get_weights() # validation loss improved so save the weights
            self.count = 0 # reset count since validation loss improved
            self.stop_count = 0
            self.best_epoch = epoch + 1 # set the value of the best epoch to this epoch
        else: # validation loss did not improve
            if self.count >= self.patience - 1: # need to adjust lr
                lr = lr * self.factor # adjust the learning rate
                self.stop_count = self.stop_count + 1 # increment stop counter because lr was adjusted
                self.count = 0 # reset counter
                tf.keras.backend.set_value(self.model.optimizer.lr, lr) # set the learning rate in the optimizer
            else:
                self.count = self.count + 1 # increment the patience counter
            if acc > self.highest_tracc:
                self.highest_tracc = acc

msg = f'{str(epoch + 1):^3s}/{str(self.epochs):4s} {loss:^9.3f}{acc * 100:^9.3f}{v_loss:^9.5f}{v_acc * 100:^9.3f}{current_lr:^9.5f}{l}
print(msg)

```

```

if self.stop_count > self.stop_patience - 1: # check if learning rate has been adjusted stop_count times with no improvement
    msg = f' training has been halted at epoch {epoch + 1} after {self.stop_patience} adjustments of learning rate with no improvement'
    print(msg)
    self.model.stop_training = True # stop training

else:
    if self.ask_epoch != None:
        if epoch + 1 >= self.ask_epoch:
            msg = 'enter H to halt training or an integer for number of epochs to run then ask again'
            print(msg)
            ans = input('')
            if ans == 'H' or ans == 'h':
                msg = f'training has been halted at epoch {epoch + 1} due to user input'
                print(msg)
                self.model.stop_training = True # stop training
            else:
                try:
                    ans = int(ans)
                    self.ask_epoch += ans
                    msg = f' training will continue until epoch ' + str(self.ask_epoch)
                    print(msg)
                    msg = '{0:^8s}{1:^10s}{2:^9s}{3:^9s}{4:^9s}{5:^9s}{6:^9s}{7:^10s}{8:^10s}{9:^8s}'.format('Epoch', 'Loss', 'Accuracy')
                    print(msg)
                except:
                    print('Invalid')

```

▼ ConvNeXtTiny

```

base_model=tf.keras.applications.ConvNeXtTiny(include_top= False, weights= "imagenet", input_shape= (224,224,3), pooling= 'max')

Downloading data from https://storage.googleapis.com/tensorflow/keras-applications/convnext/convnext\_tiny\_notop.h5
111650432/111650432 [=====] - 6s 0us/step

model = Sequential([
    base_model,
    BatchNormalization(axis= -1, momentum= 0.99, epsilon= 0.001),
    Dense(256, kernel_regularizer= regularizers.l2(l= 0.016), activity_regularizer= regularizers.l1(0.006),
          bias_regularizer= regularizers.l1(0.006), activation= 'relu'),
    Dropout(rate= 0.45, seed= 123),
    Dense(7, activation= 'softmax')
])

model.compile(Adamax(learning_rate= 0.001, weight_decay=0.02), loss=tf.keras.losses.SparseCategoricalCrossentropy(from_logits=False), metrics=)

model.summary()

Model: "sequential_1"

```

Layer (type)	Output Shape	Param #
convnext_tiny (Functional)	(None, 768)	27820128
batch_normalization (Batch Normalization)	(None, 768)	3072
dense (Dense)	(None, 256)	196864
dropout (Dropout)	(None, 256)	0
dense_1 (Dense)	(None, 7)	1799

```

Total params: 28021863 (106.89 MB)
Trainable params: 28020327 (106.89 MB)
Non-trainable params: 1536 (6.00 KB)

```

```

batch_size = 40
epochs = 40
patience = 1      # number of epochs to wait to adjust lr if monitored value does not improve
stop_patience = 3 # number of epochs to wait before stopping training if monitored value does not improve
threshold = 0.9   # if train accuracy is < threshhold adjust monitor accuracy, else monitor validation loss
factor = 0.5      # factor to reduce lr by
freeze = False    # if true free weights of the base model

```

```

ask_epoch = 5      # number of epochs to run before asking if you want to halt training
#batches = int(np.ceil(len(train_ds.labels) / batch_size))
batches = int(np.ceil(4388 / batch_size))

callbacks = [MyCallback(model= model, base_model= base_model, patience= patience,
                       stop_patience= stop_patience, threshold= threshold, factor= factor,
                       batches= batches, initial_epoch= 0, epochs= epochs, ask_epoch= ask_epoch )]

history = model.fit(x= train_ds, epochs= epochs, verbose= 0, callbacks= callbacks,
                     validation_data= val_ds, validation_steps= None, shuffle= False,
                     initial_epoch= 0)

Epoch    Loss    Accuracy   V_loss   V_acc     LR     Next LR  Monitor % Improv Duration
1 /40    7.157   43.961    6.75650  17.442   0.00100  0.00100 accuracy  0.00  1891.00
2 /40    4.994   69.690    5.07845  42.476   0.00100  0.00100 accuracy  58.53  110.89
3 /40    3.644   80.356    3.67886  59.508   0.00100  0.00100 accuracy  15.30  110.99
4 /40    2.627   87.694    2.70595  68.263   0.00100  0.00100 accuracy  9.13   110.82
5 /40    1.846   93.596    2.12952  75.718   0.00100  0.00100 val_loss  21.30  111.24
enter H to halt training or an integer for number of epochs to run then ask again
5
training will continue until epoch 10
Epoch    Loss    Accuracy   V_loss   V_acc     LR     Next LR  Monitor % Improv Duration
6 /40    1.272   97.493    1.70686  77.155   0.00100  0.00100 val_loss  19.85  110.87
7 /40    0.870   99.658    1.36678  77.497   0.00100  0.00100 val_loss  19.92  110.76
8 /40    0.618   99.932    1.19015  78.454   0.00100  0.00100 val_loss  12.92   120.46
9 /40    0.462   99.954    1.10790  78.728   0.00100  0.00100 val_loss  6.91   111.17
10 /40   0.365   100.000   0.94855  80.506   0.00100  0.00100 val_loss  14.38  110.93
enter H to halt training or an integer for number of epochs to run then ask again
2
training will continue until epoch 12
Epoch    Loss    Accuracy   V_loss   V_acc     LR     Next LR  Monitor % Improv Duration
11 /40   0.299   100.000   0.91888  79.891   0.00100  0.00100 val_loss  3.13   111.38
12 /40   0.251   100.000   0.87932  81.327   0.00100  0.00100 val_loss  4.30   111.03
enter H to halt training or an integer for number of epochs to run then ask again
2
training will continue until epoch 14
Epoch    Loss    Accuracy   V_loss   V_acc     LR     Next LR  Monitor % Improv Duration
13 /40   0.220   100.000   0.82648  81.327   0.00100  0.00100 val_loss  6.01   111.28
14 /40   0.216   99.681    0.92244  77.633   0.00100  0.00050 val_loss  -11.61  110.95
enter H to halt training or an integer for number of epochs to run then ask again
H
training has been halted at epoch 14 due to user input
training elapsed time was 1.0 hours, 2.0 minutes, 18.38 seconds)

```

```

y_test = np.concatenate([y for x, y in val_ds], axis=0)
y_pred = model.predict(val_ds)
y_pred_classes = np.argmax(y_pred, axis=1)
accuracy_score(y_test, y_pred_classes)

```

```

46/46 [=====] - 19s 259ms/step
0.8132694938440492

```

```

print(classification_report(y_test, y_pred_classes, target_names=class_names))

```

	precision	recall	f1-score	support
BAS	0.75	0.69	0.72	110
EOS	0.95	0.89	0.92	255
HAC	0.89	0.71	0.79	93
LYT	0.92	0.91	0.92	265
MON	0.75	0.85	0.80	255
NGB	0.67	0.77	0.71	235
NGS	0.80	0.72	0.76	249
accuracy			0.81	1462
macro avg	0.82	0.79	0.80	1462
weighted avg	0.82	0.81	0.81	1462

▼ ConvNeXtSmall

```

base_model=tf.keras.applications.ConvNeXtSmall(include_top= False, weights= "imagenet", input_shape= (224,224,3), pooling= 'max')

Downloading data from https://storage.googleapis.com/tensorflow/keras-applications/convnext/convnext\_small\_notop.h5
198551472/198551472 [=====] - 12s 0us/step

```

```

model = Sequential([
    base_model,
    BatchNormalization(axis= -1, momentum= 0.99, epsilon= 0.001),
    Dense(256, kernel_regularizer= regularizers.l2(l= 0.016), activity_regularizer= regularizers.l1(0.006),
          bias_regularizer= regularizers.l1(0.006), activation= 'relu'),
    Dropout(rate= 0.45, seed= 123),
    Dense(7, activation= 'softmax')
])

model.compile(Adamax(learning_rate= 0.001, weight_decay=0.02), loss=tf.keras.losses.SparseCategoricalCrossentropy(from_logits=False), metrics=)

model.summary()

Model: "sequential_2"

Layer (type)           Output Shape        Param #
=====
convnext_small (Functional (None, 768)      49454688
)
batch_normalization_1 (Batch Normalization) 3072
dense_2 (Dense)          (None, 256)       196864
dropout_1 (Dropout)       (None, 256)       0
dense_3 (Dense)          (None, 7)         1799
=====
Total params: 49656423 (189.42 MB)
Trainable params: 49654887 (189.42 MB)
Non-trainable params: 1536 (6.00 KB)

```

```

batch_size = 40
epochs = 40
patience = 1      # number of epochs to wait to adjust lr if monitored value does not improve
stop_patience = 3 # number of epochs to wait before stopping training if monitored value does not improve
threshold = 0.9   # if train accuracy is < threshold adjust monitor accuracy, else monitor validation loss
factor = 0.5      # factor to reduce lr by
freeze = False    # if true free weights of the base model
ask_epoch = 5     # number of epochs to run before asking if you want to halt training
#batches = int(np.ceil(len(train_ds.labels) / batch_size))
batches = int(np.ceil(4388 / batch_size))

callbacks = [MyCallback(model= model, base_model= base_model, patience= patience,
                       stop_patience= stop_patience, threshold= threshold, factor= factor,
                       batches= batches, initial_epoch= 0, epochs= epochs, ask_epoch= ask_epoch )]

history = model.fit(x= train_ds, epochs= epochs, verbose= 0, callbacks= callbacks,
                     validation_data= val_ds, validation_steps= None, shuffle= False,
                     initial_epoch= 0)

Epoch  Loss  Accuracy  V_loss  V_acc      LR      Next LR  Monitor % Improv Duration
1 /40  7.622  41.431  6.53380 34.063  0.00100  0.00100  accuracy  0.00  291.37
2 /40  5.280  73.108  5.13748 71.341  0.00100  0.00100  accuracy  76.46  188.71
3 /40  3.958  84.731  3.80879 78.523  0.00100  0.00100  accuracy  15.90  188.18
4 /40  2.963  91.158  2.92149 78.796  0.00100  0.00100  val_loss  23.30  188.01
5 /40  2.194  95.784  2.34063 80.643  0.00100  0.00100  val_loss  19.88  190.23
enter H to halt training or an integer for number of epochs to run then ask again
5
training will continue until epoch 10
Epoch  Loss  Accuracy  V_loss  V_acc      LR      Next LR  Monitor % Improv Duration
6 /40  1.594  98.883  2.00701 80.438  0.00100  0.00100  val_loss  14.25  188.27
7 /40  1.176  99.590  1.64974 81.259  0.00100  0.00100  val_loss  17.80  190.17
8 /40  0.869  99.863  1.48002 79.891  0.00100  0.00100  val_loss  10.29  188.34
9 /40  0.660  99.909  1.13508 84.063  0.00100  0.00100  val_loss  23.31  188.11
10 /40  0.499  100.000  0.99732 84.679  0.00100  0.00100  val_loss  12.14  188.27
enter H to halt training or an integer for number of epochs to run then ask again
5
training will continue until epoch 15
Epoch  Loss  Accuracy  V_loss  V_acc      LR      Next LR  Monitor % Improv Duration
11 /40  0.454  98.291  0.99034 80.164  0.00100  0.00100  val_loss  0.70  188.51
12 /40  0.408  97.903  1.16150 77.360  0.00100  0.00050  val_loss  -17.28  188.46
13 /40  0.316  99.521  0.93059 82.900  0.00050  0.00050  val_loss  6.03  188.28
14 /40  0.267  100.000  0.83241 83.789  0.00050  0.00050  val_loss  10.55  188.40
15 /40  0.242  100.000  0.81382 83.653  0.00050  0.00050  val_loss  2.23  188.25
enter H to halt training or an integer for number of epochs to run then ask again

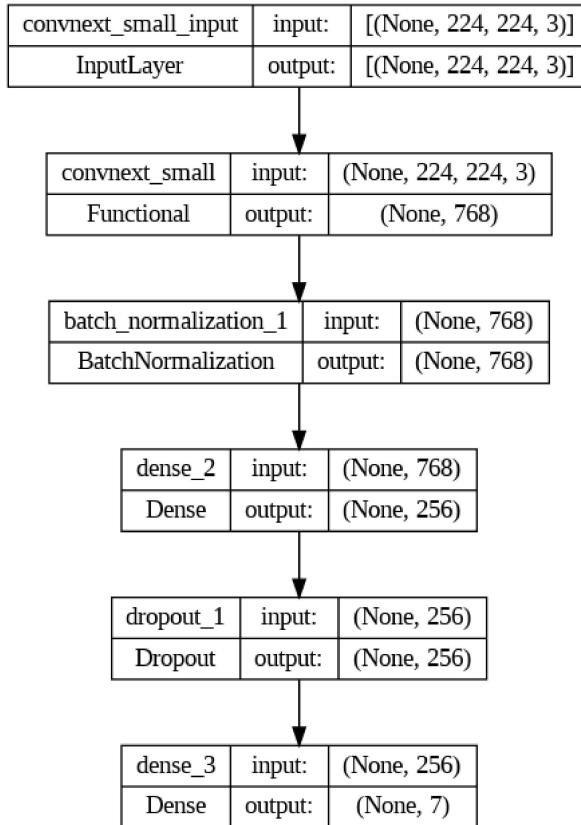
```

```

2
  training will continue until epoch 17
  Epoch    Loss   Accuracy  V_loss   V_acc     LR      Next LR  Monitor % Improv Duration
16 /40    0.225  100.000  0.81598  83.789  0.00050  0.00025  val_loss    -0.27    190.62
17 /40    0.213  100.000  0.80736  83.858  0.00025  0.00025  val_loss    0.79    188.96
enter H to halt training or an integer for number of epochs to run then ask again
1
  training will continue until epoch 18
  Epoch    Loss   Accuracy  V_loss   V_acc     LR      Next LR  Monitor % Improv Duration
18 /40    0.204  100.000  0.78687  84.542  0.00025  0.00025  val_loss    2.54    189.06
enter H to halt training or an integer for number of epochs to run then ask again
1
  training will continue until epoch 19
  Epoch    Loss   Accuracy  V_loss   V_acc     LR      Next LR  Monitor % Improv Duration
19 /40    0.199  100.000  0.80576  83.789  0.00025  0.00013  val_loss    -2.40    189.79
enter H to halt training or an integer for number of epochs to run then ask again
1
  training will continue until epoch 20
  Epoch    Loss   Accuracy  V_loss   V_acc     LR      Next LR  Monitor % Improv Duration
20 /40    0.194  100.000  0.79353  83.858  0.00013  0.00006  val_loss    -0.85    190.86
enter H to halt training or an integer for number of epochs to run then ask again
1
  training will continue until epoch 21
  Epoch    Loss   Accuracy  V_loss   V_acc     LR      Next LR  Monitor % Improv Duration
21 /40    0.191  100.000  0.78757  83.789  0.00006  0.00003  val_loss    -0.09    188.86
  training has been halted at epoch 21 after 3 adjustments of learning rate with no improvement
  training elapsed time was 1.0 hours, 18.0 minutes, 3.50 seconds)

```

```
keras.utils.plot_model(model, show_shapes=True)
```



```
#Create plots of the loss and accuracy on the training and validation sets:
acc = history.history['accuracy']
val_acc = history.history['val_accuracy']
```

```
loss = history.history['loss']
val_loss = history.history['val_loss']

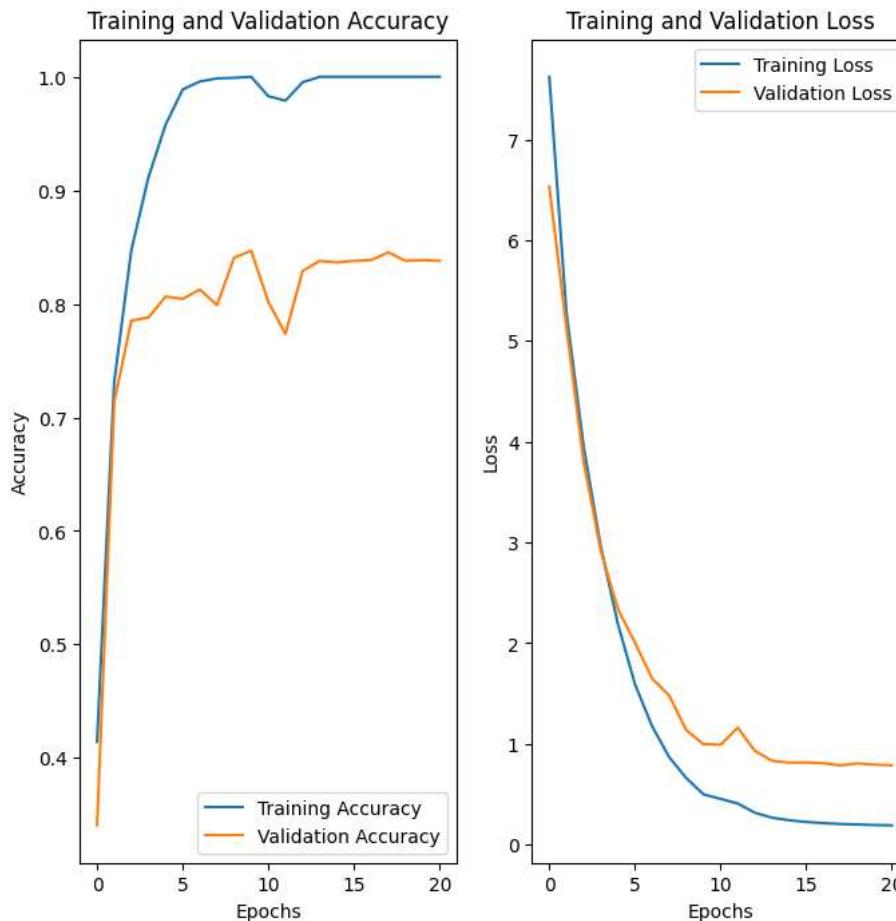
plt.figure(figsize=(8, 8))
plt.subplot(1, 2, 1)
plt.plot(acc, label='Training Accuracy')
plt.plot(val_acc, label='Validation Accuracy')
plt.legend(loc='lower right')
plt.xlabel("Epochs")
plt.ylabel("Accuracy")
```

```

plt.title('Training and Validation Accuracy')

plt.subplot(1, 2, 2)
plt.plot(loss, label='Training Loss')
plt.plot(val_loss, label='Validation Loss')
plt.legend(loc='upper right')
plt.xlabel("Epochs")
plt.ylabel("Loss")
plt.title('Training and Validation Loss')
plt.show()

```



```

y_test = np.concatenate([y for x, y in val_ds], axis=0)
y_pred = model.predict(val_ds)
y_pred_classes = np.argmax(y_pred, axis=1)
accuracy_score(y_test, y_pred_classes)

```

```

46/46 [=====] - 24s 419ms/step
0.8454172366621067

```

```
print(classification_report(y_test, y_pred_classes, target_names=class_names))
```

	precision	recall	f1-score	support
BAS	0.87	0.69	0.77	110
EOS	0.95	0.95	0.95	255
HAC	0.86	0.76	0.81	93
LYT	0.91	0.95	0.93	265
MON	0.83	0.82	0.83	255
NGB	0.76	0.78	0.77	235
NGS	0.76	0.82	0.79	249
accuracy			0.85	1462
macro avg	0.85	0.82	0.83	1462
weighted avg	0.85	0.85	0.84	1462

