

## ▼ Libraries

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import tensorflow as tf
from tensorflow import keras
from tensorflow.keras import layers

from keras.applications.densenet import DenseNet21
from keras.applications.xception import Xception
from keras.applications.densenet import preprocess_input
from keras.optimizers import Adam, SGD, Adamax
from keras.models import Model, load_model
from keras.layers import *
from sklearn.model_selection import train_test_split
from keras.callbacks import *

from tensorflow.keras.callbacks import EarlyStopping
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, Embedding, LSTM, Conv1D, MaxPool1D
from sklearn.model_selection import train_test_split
from sklearn.metrics import classification_report, accuracy_score
from tensorflow.keras import regularizers
from keras.src import callbacks
```

## ▼ Train test split

```
from google.colab import drive
drive.mount('/content/drive', force_remount=True)

Mounted at /content/drive

image_size = (224, 224)
batch_size = 32

train_ds, val_ds = tf.keras.utils.image_dataset_from_directory(
    "/content/drive/MyDrive/BM1000",
    validation_split=0.25,
    subset="both",
    seed=1337,
    image_size=image_size,
    batch_size=batch_size,
)
```

Found 5850 files belonging to 7 classes.  
Using 4388 files for training.  
Using 1462 files for validation.

```
train_ds

<_PrefetchDataset element_spec=(TensorSpec(shape=(None, 224, 224, 3), dtype=tf.float32, name=None), TensorSpec(shape=(None,), dtype=tf.int32, name=None))>

class_names = train_ds.class_names
print(class_names)
num_classes = len(class_names)
print(num_classes)


['BAS', 'EOS', 'HAC', 'LYT', 'MON', 'NGB', 'NGS']

dataset_unbatched = tuple(train_ds.unbatch())
labels = []
for (image,label) in dataset_unbatched:
    labels.append(label.numpy())
labels = pd.Series(labels)

# adjustments
count = labels.value_counts().sort_index()
#count.index
print(count)

0    331
1    745
2    316
3    735
4    745
5    765
6    751
dtype: int64
```

Here are the first nine images from the training dataset.  
plt.figure(figsize=(10, 10))  
for images, labels in train\_ds.take(1):  
 for i in range(9):  
 ax = plt.subplot(3, 3, i + 1)  
 plt.imshow(images[i].numpy().astype("uint8"))  
 plt.title(class\_names[labels[i]])  
 plt.axis("off")



```
for image_batch, labels_batch in train_ds:  
    print(image_batch.shape)  
    print(labels_batch.shape)  
    break
```

(32, 224, 224, 3)  
(32,)

```
image_batch, labels_batch = next(iter(train_ds))
first_image = image_batch[0]
print(first_image)
```

```
tf.Tensor([[[[102., 44., 224.]
            [102., 47., 228.]
            [ 95., 41., 225.]
            ...
            [189., 179., 239.]
            [199., 188., 246.]
            [197., 186., 244.]]]

[[105., 47., 227.]
            [102., 47., 228.]
            [ 94., 40., 222.]
            ...
            [187., 177., 237.]
            [193., 183., 243.]
            [191., 181., 241.]]]

[[103., 45., 227.]])
```

```
[101. 46. 227.]
[ 96. 43. 223.]
...
[199. 190. 253.]
[206. 197. 255.]
[205. 196. 255.]]
```

...

```
[[220. 214. 252.]
[219. 213. 251.]
[221. 213. 254.]
...
[208. 196. 246.]
[204. 191. 244.]
[206. 193. 247.]]]
```

```
[[221. 217. 254.]
[218. 212. 250.]
[221. 213. 252.]
...
[195. 183. 231.]
[205. 191. 242.]
[206. 189. 241.]]]
```


```
[[217. 213. 250.]
[220. 216. 253.]
[220. 212. 251.]
...
[202. 190. 236.]
[205. 189. 238.]
[228. 210. 255.]]], shape=(224, 224, 3), dtype=float32)
```

## ▼ Data Augmentation

```
data_augmentation = keras.Sequential(
[
    layers.RandomFlip("horizontal",
        input_shape=(224,
                    224,
                    3)),
    layers.RandomRotation(0.2),
    layers.RandomZoom(0.2),
    layers.RandomContrast(0.2),
    layers.RandomBrightness(0.2)
])
```

#Visualize a few augmented examples by applying data augmentation to the same image several times:

```
plt.figure(figsize=(10, 10))
for images, _ in train_ds.take(1):
    for i in range(9):
        augmented_images = data_augmentation(images)
        ax = plt.subplot(3, 3, i + 1)
        plt.imshow(augmented_images[0].numpy().astype("uint8"))
        plt.axis("off")
```



```
train_ds = train_ds.map(
    lambda x, y: (data_augmentation(x, training=True), y))
```

## ▼ Configure the dataset for performance

Make sure to use buffered prefetching, so you can yield data from disk without having I/O become blocking. These are two important methods you should use when loading data:

Dataset.cache keeps the images in memory after they're loaded off disk during the first epoch. This will ensure the dataset does not become a bottleneck while training your model. If your dataset is too large to fit into memory, you can also use this method to create a performant on-disk cache.

Dataset.prefetch overlaps data preprocessing and model execution while training.

```
AUTOTUNE = tf.data.AUTOTUNE
train_ds = train_ds.cache().shuffle(1000).prefetch(buffer_size=AUTOTUNE)
val_ds = val_ds.cache().prefetch(buffer_size=AUTOTUNE)
```

## ▼ Early Stopping

```
# Define EarlyStopping callback
early_stopping = EarlyStopping(monitor='val_loss', patience=10, restore_best_weights=True)
```

## ▼ Sequential Model 1

```
model = Sequential([
    layers.Rescaling(1./255, input_shape=(224, 224, 3)),
    layers.Conv2D(16, 3, padding='same', activation='relu'),
    layers.MaxPooling2D(),
    layers.Conv2D(32, 3, padding='same', activation='relu'),
    layers.MaxPooling2D(),
    layers.Conv2D(64, 3, padding='same', activation='relu', kernel_regularizer=regularizers.l2(l=0.01)),
    layers.MaxPooling2D(),
    layers.Dropout(0.5),
    layers.Flatten(),
    layers.Dense(128, activation='relu', kernel_regularizer=regularizers.l2(l=0.01)),
    layers.Dense(7, activation='softmax', name="outputs")
])
```

```
model.compile(Adamax(learning_rate= 0.001), loss=tf.keras.losses.SparseCategoricalCrossentropy(from_logits=False), metrics= ['accuracy'])
```

```
model.summary()
```

```
Model: "sequential_1"
-----
```


Layer (type)	Output Shape	Param #
rescaling (Rescaling)	(None, 224, 224, 3)	0
conv2d (Conv2D)	(None, 224, 224, 16)	448
max_pooling2d (MaxPooling2D)	(None, 112, 112, 16)	0
conv2d_1 (Conv2D)	(None, 112, 112, 32)	4640
max_pooling2d_1 (MaxPooling2D)	(None, 56, 56, 32)	0
conv2d_2 (Conv2D)	(None, 56, 56, 64)	18496
max_pooling2d_2 (MaxPooling2D)	(None, 28, 28, 64)	0

```

dropout (Dropout)      (None, 28, 28, 64)      0
flatten (Flatten)     (None, 50176)            0
dense (Dense)         (None, 128)              6422656
outputs (Dense)       (None, 7)                903
=====
Total params: 6447143 (24.59 MB)
Trainable params: 6447143 (24.59 MB)
Non-trainable params: 0 (0.00 Byte)

```

```
keras.utils.plot_model(model, show_shapes=True)
```



```

epochs = 70
history = model.fit(
    train_ds,
    validation_data=val_ds,
    epochs=epochs,
    batch_size=256,
    callbacks=[early_stopping]
)

Epoch 1/70
138/138 [=====] - 95s 216ms/step - loss: 3.0430 - accuracy: 0.2359 - val_loss: 2.4720 - val_accuracy: 0.3495
Epoch 2/70
138/138 [=====] - 6s 45ms/step - loss: 2.3499 - accuracy: 0.3336 - val_loss: 2.1710 - val_accuracy: 0.4001
Epoch 3/70
138/138 [=====] - 6s 45ms/step - loss: 2.1544 - accuracy: 0.3546 - val_loss: 2.0539 - val_accuracy: 0.4049
Epoch 4/70
138/138 [=====] - 6s 46ms/step - loss: 2.0373 - accuracy: 0.3751 - val_loss: 2.0113 - val_accuracy: 0.3495
Epoch 5/70
138/138 [=====] - 6s 45ms/step - loss: 1.9520 - accuracy: 0.3838 - val_loss: 1.8541 - val_accuracy: 0.4220
Epoch 6/70
138/138 [=====] - 6s 45ms/step - loss: 1.8773 - accuracy: 0.3906 - val_loss: 1.8111 - val_accuracy: 0.4337
Epoch 7/70
138/138 [=====] - 6s 46ms/step - loss: 1.8185 - accuracy: 0.4191 - val_loss: 1.7925 - val_accuracy: 0.4193
Epoch 8/70
138/138 [=====] - 6s 45ms/step - loss: 1.7857 - accuracy: 0.4389 - val_loss: 1.7400 - val_accuracy: 0.4316
Epoch 9/70
138/138 [=====] - 6s 46ms/step - loss: 1.7280 - accuracy: 0.4464 - val_loss: 1.6406 - val_accuracy: 0.4836
Epoch 10/70
138/138 [=====] - 6s 45ms/step - loss: 1.6919 - accuracy: 0.4505 - val_loss: 1.6554 - val_accuracy: 0.4740
Epoch 11/70
138/138 [=====] - 7s 48ms/step - loss: 1.6459 - accuracy: 0.4756 - val_loss: 1.6161 - val_accuracy: 0.4822
Epoch 12/70
138/138 [=====] - 6s 47ms/step - loss: 1.5792 - accuracy: 0.5021 - val_loss: 1.5628 - val_accuracy: 0.5021
Epoch 13/70
138/138 [=====] - 6s 46ms/step - loss: 1.5339 - accuracy: 0.5312 - val_loss: 1.5778 - val_accuracy: 0.4850
Epoch 14/70
138/138 [=====] - 6s 47ms/step - loss: 1.4890 - accuracy: 0.5558 - val_loss: 1.5733 - val_accuracy: 0.5048
Epoch 15/70
138/138 [=====] - 6s 47ms/step - loss: 1.4568 - accuracy: 0.5675 - val_loss: 1.6176 - val_accuracy: 0.4973
Epoch 16/70
138/138 [=====] - 7s 47ms/step - loss: 1.4197 - accuracy: 0.5932 - val_loss: 1.5258 - val_accuracy: 0.5321
Epoch 17/70
138/138 [=====] - 6s 46ms/step - loss: 1.3787 - accuracy: 0.6051 - val_loss: 1.5325 - val_accuracy: 0.5417
Epoch 18/70
138/138 [=====] - 7s 48ms/step - loss: 1.3363 - accuracy: 0.6335 - val_loss: 1.5519 - val_accuracy: 0.5308
Epoch 19/70
138/138 [=====] - 6s 47ms/step - loss: 1.2957 - accuracy: 0.6522 - val_loss: 1.5460 - val_accuracy: 0.5472
Epoch 20/70
138/138 [=====] - 7s 47ms/step - loss: 1.2407 - accuracy: 0.6718 - val_loss: 1.4924 - val_accuracy: 0.5800
Epoch 21/70
138/138 [=====] - 7s 47ms/step - loss: 1.2095 - accuracy: 0.6823 - val_loss: 1.5973 - val_accuracy: 0.5595
Epoch 22/70
138/138 [=====] - 6s 47ms/step - loss: 1.1600 - accuracy: 0.7053 - val_loss: 1.5541 - val_accuracy: 0.5636
Epoch 23/70
138/138 [=====] - 7s 48ms/step - loss: 1.1397 - accuracy: 0.7217 - val_loss: 1.6633 - val_accuracy: 0.5404
Epoch 24/70
138/138 [=====] - 6s 47ms/step - loss: 1.0961 - accuracy: 0.7297 - val_loss: 1.5710 - val_accuracy: 0.5711
Epoch 25/70
138/138 [=====] - 7s 48ms/step - loss: 1.0514 - accuracy: 0.7523 - val_loss: 1.6110 - val_accuracy: 0.5698
Epoch 26/70
138/138 [=====] - 7s 47ms/step - loss: 1.0366 - accuracy: 0.7596 - val_loss: 1.6294 - val_accuracy: 0.5554
Epoch 27/70
138/138 [=====] - 7s 48ms/step - loss: 1.0002 - accuracy: 0.7755 - val_loss: 1.6467 - val_accuracy: 0.5527
Epoch 28/70
138/138 [=====] - 7s 48ms/step - loss: 0.9601 - accuracy: 0.7892 - val_loss: 1.6799 - val_accuracy: 0.5458
Epoch 29/70
138/138 [=====] - 7s 49ms/step - loss: 0.9578 - accuracy: 0.7890 - val_loss: 1.7277 - val_accuracy: 0.5568

```

```
#Create plots of the loss and accuracy on the training and validation sets:
acc = history.history['accuracy']
val_acc = history.history['val_accuracy']
```

```
loss = history.history['loss']
val_loss = history.history['val_loss']
```

```
plt.figure(figsize=(8, 8))
plt.subplot(1, 2, 1)
plt.plot(acc, label='Training Accuracy')
plt.plot(val_acc, label='Validation Accuracy')
plt.legend(loc='lower right')
plt.xlabel("Epochs")
plt.ylabel("Accuracy")
plt.title('Training and Validation Accuracy')

plt.subplot(1, 2, 2)
plt.plot(loss, label='Training Loss')
plt.plot(val_loss, label='Validation Loss')
plt.legend(loc='upper right')
plt.xlabel("Epochs")
plt.ylabel("Loss")
plt.title('Training and Validation Loss')
plt.show()
```



```
y_test = np.concatenate([y for x, y in val_ds], axis=0)
y_pred = model.predict(val_ds)
y_pred_classes = np.argmax(y_pred, axis=1)
accuracy_score(y_test, y_pred_classes)
```

```
46/46 [=====] - 1s 11ms/step
```

```
0.5800273597811217
```

```
print(classification_report(y_test, y_pred_classes, target_names=class_names))
```

	precision	recall	f1-score	support
BAS	0.44	0.16	0.24	110
EOS	0.60	0.67	0.63	255
HAC	0.60	0.56	0.58	93
LYT	0.69	0.74	0.71	265
MON	0.53	0.63	0.57	255
NGB	0.51	0.41	0.45	235
NGS	0.57	0.62	0.59	249
<b>accuracy</b>	<b>0.56</b>	<b>0.54</b>	<b>0.58</b>	<b>1462</b>
<b>macro avg</b>	<b>0.56</b>	<b>0.54</b>	<b>0.54</b>	<b>1462</b>
<b>weighted avg</b>	<b>0.57</b>	<b>0.58</b>	<b>0.57</b>	<b>1462</b>